

Министерство образования Республики Беларусь
Учреждение образования
«Брестский Государственный технический университет»
Кафедра ИИТ

Отчет по лабораторной работе 3

Специальность ИИ-23

Выполнил:

Макаревич Н.Р.

Студент группы ИИ-23

Проверил:

Андренко К. В.

Преподаватель-стажёр
Кафедры ИИТ,

«___» _____ 2025 г.

Цель: научиться осуществлять предобучение нейронных сетей с помощью автоэнкодерного подхода

Общее задание

1. Взять за основу любую сверточную или полносвязную архитектуру с количеством слоев более 3. Осуществить ее обучение (без предобучения) в соответствии с вариантом задания. Получить оценку эффективности модели, используя метрики, специфичные для решаемой задачи (например, MAPE – для регрессионной задачи или F1/Confusion matrix для классификационной).
2. Выполнить обучение с предобучением, используя автоэнкодерный подход, алгоритм которого изложен в лекции. Условие останова (например, по количеству эпох) при обучении отдельных слоев с использованием автоэнкодера выбрать самостоятельно.
3. Сравнить результаты, полученные при обучении с/без предобучения, сделать выводы.
4. Выполните пункты 1-3 для датасетов из ЛР 2.
5. Оформить отчет по выполненной работе, загрузить исходный код и отчет в соответствующий репозиторий на github.

Код программы:

```
# -*- coding: utf-8 -*-
```

```
"""iad_3.ipynb
```

Automatically generated by Colab.

Original file is located at

https://colab.research.google.com/drive/1NbHINh_nwV_Emt9h_uLQDu0YH4HbFS5x

```
"""
```

```
!pip install ucimlrepo --quiet
```

```
!pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cpu -q
```

!pip install scikit-learn pandas numpy matplotlib seaborn --quiet

```
import random
```

```
import numpy as np
```

```
import pandas as pd
```

```
import torch
```

```
import torch.nn as nn
```

```
import torch.optim as optim
```

```
from torch.utils.data import TensorDataset, DataLoader
```

```
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import classification_report, confusion_matrix, f1_score,  
accuracy_score
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
RND = 42
```

```
random.seed(RND)
```

```
np.random.seed(RND)
```

```
torch.manual_seed(RND)
```

```
device = torch.device("cpu")
```

```
from ucimlrepo import fetch_ucirepo
```

```
ds = fetch_ucirepo(id=863)
```

```
X = ds.data.features
```

```
y = ds.data.targets
```

```
if not isinstance(X, pd.DataFrame):
```

```
    X = pd.DataFrame(X)
```

```
if isinstance(y, (pd.DataFrame, pd.Series)):
```

```
    y = np.array(y).reshape(-1,)
```

```
else:
```

```
    y = np.array(y)
```

```
le = LabelEncoder()
```

```
y_enc = le.fit_transform(y)
```

```
class_names = le.classes_
```

```
n_classes = len(class_names)
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X.values.astype(float))
```

```
X_train, X_temp, y_train, y_temp = train_test_split(
```

```
    X_scaled, y_enc, test_size=0.4, random_state=RND, stratify=y_enc
```

```
)
```

```
X_val, X_test, y_val, y_test = train_test_split(
```

```
    X_temp, y_temp, test_size=0.5, random_state=RND, stratify=y_temp
```

```
)
```

```
def to_loader(X, y, batch_size=32, shuffle=True):
```

```
    X_t = torch.tensor(X, dtype=torch.float32)
```

```
    y_t = torch.tensor(y, dtype=torch.long)
```

```

ds = TensorDataset(X_t, y_t)

return DataLoader(ds, batch_size=batch_size, shuffle=shuffle)

train_loader = to_loader(X_train, y_train, batch_size=32, shuffle=True)
val_loader = to_loader(X_val, y_val, batch_size=64, shuffle=False)
test_loader = to_loader(X_test, y_test, batch_size=64, shuffle=False)

def evaluate_model(model, loader):
    model.eval()

    ys = []
    ys_pred = []

    with torch.no_grad():
        for xb, yb in loader:
            xb = xb.to(device)
            out = model(xb)
            preds = out.argmax(dim=1).cpu().numpy()
            ys_pred.extend(preds.tolist())
            ys.extend(yb.numpy().tolist())

    return np.array(ys), np.array(ys_pred)

def print_metrics(y_true, y_pred, labels=None):
    print("Accuracy:", accuracy_score(y_true, y_pred))
    print("Macro F1: ", f1_score(y_true, y_pred, average='macro'))
    print("\nClassification report:\n", classification_report(y_true, y_pred,
target_names=labels))

    cm = confusion_matrix(y_true, y_pred)

    plt.figure(figsize=(6,5))

```

```
sns.heatmap(cm, annot=True, fmt='d', xticklabels=labels, yticklabels=labels,
cmap='Blues')
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("True")
```

```
plt.show()
```

```
input_dim = X_train.shape[1]
```

```
hidden_sizes = [64, 32, 16]
```

```
dropout = 0.2
```

```
class FCNet(nn.Module):
```

```
    def __init__(self, input_dim, hidden_sizes, n_classes, dropout=0.2):
```

```
        super().__init__()
```

```
        layers = []
```

```
        prev = input_dim
```

```
        for h in hidden_sizes:
```

```
            layers.append(nn.Linear(prev, h))
```

```
            layers.append(nn.ReLU())
```

```
            layers.append(nn.Dropout(dropout))
```

```
            prev = h
```

```
        layers.append(nn.Linear(prev, n_classes))
```

```
        self.net = nn.Sequential(*layers)
```

```
    def forward(self, x):
```

```
        return self.net(x)
```

```
model_plain = FCNet(input_dim, hidden_sizes, n_classes, dropout=dropout).to(device)
```

```
print(model_plain)
```

```
def train_classifier(model, train_loader, val_loader, epochs=100, lr=1e-3, weight_decay=1e-5, patience=10):
```

```
    criterion = nn.CrossEntropyLoss()
```

```
    optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)
```

```
    best_val_loss = float('inf')
```

```
    best_state = None
```

```
    patience_cnt = 0
```

```
    history = {'train_loss':[], 'val_loss':[]}
```

```
    for ep in range(1, epochs+1):
```

```
        model.train()
```

```
        total_loss = 0.0
```

```
        for xb, yb in train_loader:
```

```
            xb, yb = xb.to(device), yb.to(device)
```

```
            optimizer.zero_grad()
```

```
            out = model(xb)
```

```
            loss = criterion(out, yb)
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
            total_loss += loss.item() * xb.size(0)
```

```
    avg_train_loss = total_loss / len(train_loader.dataset)
```

```
    model.eval()
```

```
    total_val_loss = 0.0
```

```
    with torch.no_grad():
```

```
        for xb, yb in val_loader:
```

```
            xb, yb = xb.to(device), yb.to(device)
```

```
            out = model(xb)
```

```

    loss = criterion(out, yb)

    total_val_loss += loss.item() * xb.size(0)

avg_val_loss = total_val_loss / len(val_loader.dataset)

history['train_loss'].append(avg_train_loss)

history['val_loss'].append(avg_val_loss)


if avg_val_loss < best_val_loss - 1e-5:

    best_val_loss = avg_val_loss

    best_state = {k:v.cpu() for k,v in model.state_dict().items()}

    patience_cnt = 0

else:

    patience_cnt += 1

if patience_cnt >= patience:

    print(f"Early stopping at epoch {ep}")

    break

if ep % 10 == 0 or ep==1:

    print(f"Epoch {ep} train_loss={avg_train_loss:.4f} val_loss={avg_val_loss:.4f}")


if best_state is not None:

    model.load_state_dict(best_state)

return model, history

```

```

model_plain_trained, hist_plain = train_classifier(model_plain, train_loader, val_loader,
epochs=200, lr=1e-3, patience=15)

```

```

y_true, y_pred = evaluate_model(model_plain_trained, test_loader)

print_metrics(y_true, y_pred, labels=class_names)

```


ae_epochs = 80

ae_lr = 1e-3

ae_batch = 32

```
class SimpleAE(nn.Module):
```

```
    def __init__(self, input_dim, bottleneck_dim):
```

```
        super().__init__()
```

```
        self.enc = nn.Sequential(nn.Linear(input_dim, bottleneck_dim), nn.ReLU())
```

```
        self.dec = nn.Sequential(nn.Linear(bottleneck_dim, input_dim))
```

```
    def forward(self, x):
```

```
        z = self.enc(x)
```

```
        xr = self.dec(z)
```

```
        return xr, z
```

```
def train_autoencoder(ae, data_X, epochs=50, batch_size=32, lr=1e-3, verbose=False):
```

```
    ae = ae.to(device)
```

```
    opt = optim.Adam(ae.parameters(), lr=lr)
```

```
    criterion = nn.MSELoss()
```

```
    loader = to_loader(data_X, np.zeros(len(data_X)), batch_size=batch_size, shuffle=True)
```

```
    for ep in range(1, epochs+1):
```

```
        ae.train()
```

```
        total = 0.0
```

```
        for xb, _ in loader:
```

```
            xb = xb.to(device)
```

```
            xr, _ = ae(xb)
```

```

    loss = criterion(xr, xb)

    opt.zero_grad()

    loss.backward()

    opt.step()

    total += loss.item() * xb.size(0)

    if verbose and (ep==1 or ep%20==0 or ep==epochs):
        print(f"AE epoch {ep} loss {total/len(data_X):.6f}")

    return ae

ae_models = []

activations_train = X_train.copy()
activations_val = X_val.copy()
activations_test = X_test.copy()

trained_encoders = []

for i, h in enumerate(hidden_sizes):

    ae = SimpleAE(input_dim=activations_train.shape[1], bottleneck_dim=h)

    ae = train_autoencoder(ae, activations_train, epochs=ae_epochs, batch_size=ae_batch,
lr=ae_lr, verbose=True)

    ae_models.append(ae)

    ae = ae.to(device)

    ae.eval()

    with torch.no_grad():

        train_z = ae.enc(torch.tensor(activations_train,
dtype=torch.float32).to(device)).cpu().numpy()

```

```
val_z = ae.enc(torch.tensor(activations_val,
dtype=torch.float32).to(device)).cpu().numpy()
```

```
test_z = ae.enc(torch.tensor(activations_test,
dtype=torch.float32).to(device)).cpu().numpy()
```

```
linear_layer = ae.enc[0]
```

```
enc_linear = nn.Linear(linear_layer.in_features, linear_layer.out_features)
```

```
enc_linear.weight.data = linear_layer.weight.data.cpu().clone()
```

```
enc_linear.bias.data = linear_layer.bias.data.cpu().clone()
```

```
trained_encoders.append(nn.Sequential(enc_linear, nn.ReLU()))
```

```
activations_train = train_z
```

```
activations_val = val_z
```

```
activations_test = test_z
```

```
print("\nPretraining finished. Number of pretrained encoders:", len(trained_encoders))
```

```
class FCNetFromEncoders(nn.Module):
```

```
    def __init__(self, encoders, n_classes, dropout=0.2):
```

```
        super().__init__()
```

```
        layers = []
```

```
        for enc in encoders:
```

```
            layers.append(enc[0])
```

```
            layers.append(nn.ReLU())
```

```
            layers.append(nn.Dropout(dropout))
```

```

last_dim = encoders[-1][0].out_features

layers.append(nn.Linear(last_dim, n_classes))

self.net = nn.Sequential(*layers)

def forward(self, x):

    return self.net(x)

model_pretrained = FCNetFromEncoders(trained_encoders, n_classes,
dropout=dropout).to(device)

print(model_pretrained)

model_pt_trained, hist_pt = train_classifier(model_pretrained, train_loader, val_loader,
epochs=200, lr=1e-3, patience=15)

y_true_pt, y_pred_pt = evaluate_model(model_pt_trained, test_loader)

print_metrics(y_true_pt, y_pred_pt, labels=class_names)

print("Baseline (no pretraining):")

print("Accuracy:", accuracy_score(y_true, y_pred))

print("Macro F1:", f1_score(y_true, y_pred, average='macro'))

print()



print("With pretraining:")

print("Accuracy:", accuracy_score(y_true_pt, y_pred_pt))

print("Macro F1:", f1_score(y_true_pt, y_pred_pt, average='macro'))

```


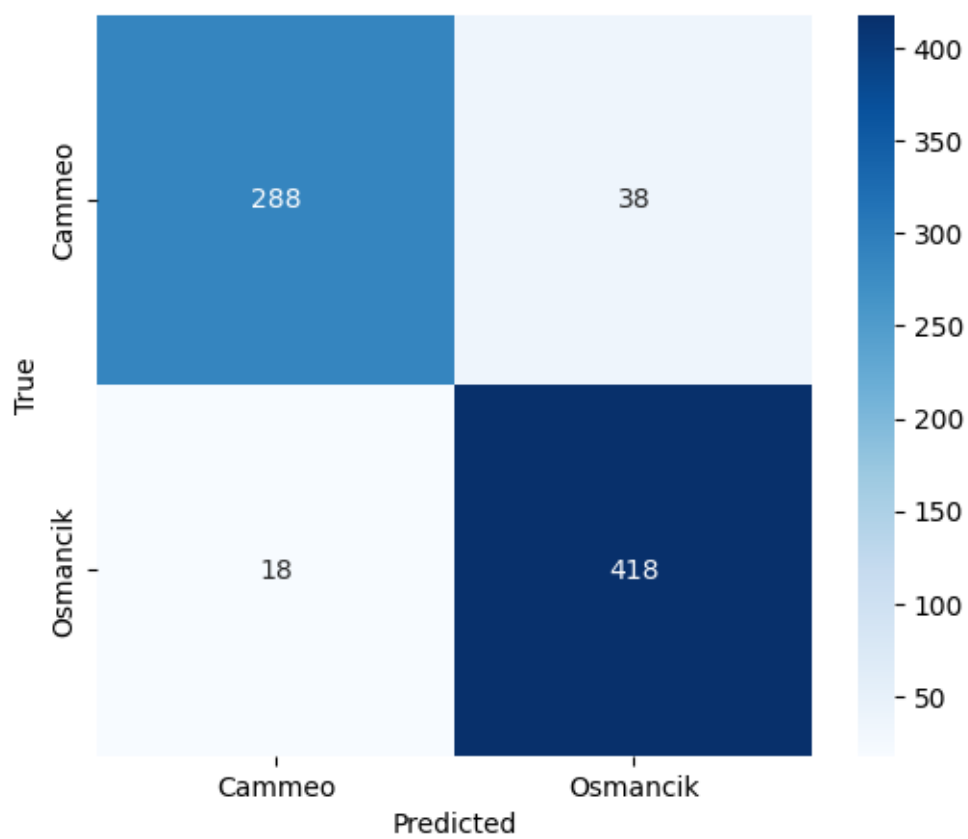
Датасет 3 лаб работы:



Baseline (no pretraining):
Accuracy: 0.729064039408867
Macro F1: 0.7270163798465684

With pretraining:
Accuracy: 0.6995073891625616
Macro F1: 0.6854236536016316

Датасет 2 лаб работы:



```
Baseline (no pretraining):  
Accuracy: 0.926509186351706  
Macro F1: 0.9243060680024976
```

```
With pretraining:  
Accuracy: 0.9304461942257218  
Macro F1: 0.9283939979042665
```

Вывод: научился предобучать НС с помощью автоенкодеров.